


Exercise: Monte Carlo Methods and Temporal-Difference Learning

Task 1)


- What is the difference between Monte Carlo (MC) methods and Dynamic Programming?
- What is bootstrapping? Is MC using bootstrapping?
- What is the difference between on-policy and off-policy?
- What is the following equation describing?

$$\rho_{t:T-1} \doteq \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}$$

Task 2) Frozen Lake



MC -> Policy in Bild
eingezeichnet nach 2000
Episoden und Epsilon-
Greedy Sequenz
Generation



Q-Learning -> Policy in
Bild eingezeichnet nach
20000 Episoden und
Epsilon-Greedy Sequenz
Generation

Solve the Open AI task frozen lake (desc=None,map_name="4x4", is_slippery=False) using MC and Q-Learning. You might either start your own implementation or complete the Python script provided in the lecture.

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

```
Initialize:
   $\pi(s) \in \mathcal{A}(s)$  (arbitrarily), for all  $s \in \mathcal{S}$ 
   $Q(s, a) \in \mathbb{R}$  (arbitrarily), for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
   $Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 

Loop forever (for each episode):
  Choose  $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$  randomly such that all pairs have probability  $> 0$ 
  Generate an episode from  $S_0, A_0$ , following  $\pi$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ 
   $G \leftarrow 0$ 
  Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :
     $G \leftarrow \gamma G + R_{t+1}$ 
    Unless the pair  $S_t, A_t$  appears in  $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$ :
      Append  $G$  to  $Returns(S_t, A_t)$ 
     $Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$ 
     $\pi(S_t) \leftarrow \arg\max_a Q(S_t, a)$ 
```

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

```
Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$ 
Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$ 

Loop for each episode:
  Initialize  $S$ 
  Loop for each step of episode:
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)
    Take action  $A$ , observe  $R, S'$ 
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```

Using Open AI Gym's Frozen Lake Environment

https://www.gymnasium.ml/environments/toy_text/frozen_lake/

> pip install gym # to install gym lib

> pip install pygame # to install pygame lib (needed for visualization)

Suggested further reading:

<https://blog.paperspace.com/getting-started-with-openai-gym/>

<https://www.learndatasci.com/tutorials/reinforcement-q-learning-scratch-python-openai-gym/>

Bonus: solve the task using is_slippery=True

Task 3) Cliff Walk

Consider the gridworld shown below. This is a standard undiscounted, episodic task, with start and goal states, and the usual actions causing movement up, down, right, and left. Reward is -1 on all transitions except those into the region marked “The Cliff”. Stepping into this region incurs a reward of -100 and sends the agent instantly back to the start. Movements outside the gridworld result in a reward of -1 and the state remains the same.

Use Sarsa or Q-learning methods with ϵ -greedy action selection ($\epsilon = 0.1$) for 500 episodes. Draw the final walk of the agent in episode 500. Does it resemble the path shown below?

